

STA 314: Statistical Methods for Machine Learning I

Lecture - k-nearest neighbour (classification)

Xin Bing

Department of Statistical Sciences
University of Toronto

A Classical Local Approach: Nearest Neighbors

- Suppose we're given a new feature vector $\mathbf{x} \in \mathbb{R}^p$ we consider classification.
- The idea: find the nearest feature vector to \mathbf{x} in the training set and use its label.
- Can formalize "nearest" in terms of the Euclidean distance

$$\|\mathbf{x}_i - \mathbf{x}_{i'}\|_2 = \sqrt{\sum_{j=1}^p (x_{ij} - x_{i'j})^2}$$

Algorithm (1-NN):


1. Find example (\mathbf{x}_*, y_*) (from the stored training data) closest to \mathbf{x} . That is:

$$\mathbf{x}_* = \underset{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}}{\operatorname{argmin}} \text{distance}(\mathbf{x}_i, \mathbf{x})$$

2. Output y_* as the label


Nearest Neighbors: Decision Boundaries

We can visualize the behavior in the classification setting using a **Voronoi diagram**.




Nearest Neighbors: Decision Boundaries

Decision boundary: the boundary between regions of the feature space assigned to different categories.



Nearest Neighbors: Decision Boundaries



Example: 2D decision boundary

k -Nearest Neighbors


Algorithm (k -NN):

1. Find k data points $(\mathbf{x}_{(1)}, y_{(1)}), \dots, (\mathbf{x}_{(k)}, y_{(k)})$ closest to the test instance \mathbf{x}
2. Classification output is majority class

$$\arg \max_{y \in C} \sum_{i=1}^k \mathbb{1}\{y = y_{(i)}\}.$$


k -NN

$k=1$



k -NN

$k=15$



k -NN

Tradeoffs in choosing k

- Small k
 - ▶ More flexible decision boundary but high variance
 - ▶ Overfitting: sensitive to random noise in the training data.
- Large k
 - ▶ Less flexible decision boundary and smaller variance
 - ▶ Underfitting: fail to capture the true decision boundary.
- Balancing k
 - ▶ Optimal choice of k depends on number of data points n .
 - ▶ Nice theoretical properties if

$$k \rightarrow \infty, \quad \text{and} \quad \frac{k}{n} \rightarrow 0 \quad (\text{ESL 2.4}).$$

- ▶ Rule of thumb: choose k from $[1, \sqrt{n}]$ via cross-validation!

Pitfalls: The Curse of Dimensionality

- k -NN suffers the curse of dimensionality!
 - ▶ In high dimensions, “most” points are approximately the same distance because they are far away from each other.
- Saving grace: some datasets (e.g. images) may have low **intrinsic dimension**, i.e. lie on or near a low-dimensional manifold.





Image credit: https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_swiss_roll.html

The neighborhood structure depends on the intrinsic dimension.

Pitfalls: Normalization

- Nearest neighbors can be sensitive to the ranges of different features.
- Often, the units are arbitrary:




- Simple fix: **standardize** each dimension to be zero mean and unit variance.
- Caution: depending on the problem, the scale might be important!

Pitfalls: Computational Cost

- Computational cost for **training**: 0
- Computational cost for classifying **test data**, per data point (un-modified algorithm)
 - ▶ Calculate p -dimensional Euclidean distances with n data points:
 $\mathcal{O}(np)$
 - ▶ Sort the distances: $\mathcal{O}(n \log n)$
- This must be done for *each* test data point, which is very expensive by the standards of a learning algorithm!
- Need to store the entire dataset in memory!
- Tons of work has gone into algorithms and data structures for efficient nearest neighbors with high dimensions and/or large datasets.

Example: Digit Classification

- Decent performance when lots of data




- Yann LeCunn – MNIST Digit Recognition
 - Handwritten digits
 - 28x28 pixel images: $d = 784$
 - 60,000 training samples
 - 10,000 test samples
- Nearest neighbour is competitive

	Test Error Rate (%)
Linear classifier (1-layer NN)	12.0
K-nearest-neighbors, Euclidean	5.0
K-nearest-neighbors, Euclidean, deskewed	2.4
K-NN, Tangent Distance, 16x16	1.1
K-NN, shape context matching	0.67
1000 RBF + linear classifier	3.6
SVM deg 4 polynomial	1.1
2-layer NN, 300 hidden units	4.7
2-layer NN, 300 HU, [deskewing]	1.6
LeNet-5, [distortions]	0.8
Boosted LeNet-4, [distortions]	0.7

Example: Digit Classification

- **Changing the distance measure** can really improve k -NN.
- Example: shape contexts for object recognition. In order to achieve invariance to image transformations, they tried to warp one image to match the other image.
 - ▶ Distance measure: average distance between corresponding points on *warped* images
- Achieved 0.63% error on MNIST, compared with 3% for Euclidean KNN.
- Competitive with the state of the art at the time, but required careful engineering.



[Belongie, Malik, and Puzicha, 2002. Shape matching and object recognition using shape contexts.]